

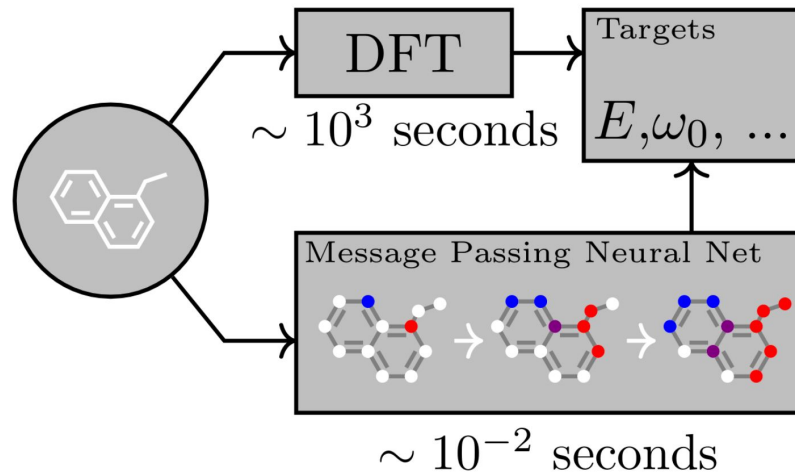
Neural Message Passing for Quantum Chemistry

Motivation

- Applications of quantum chemistry:
 - ➡ Drug design and discovery
 - ➡ Toxicity analysis
 - ➡ Material science
 - ➡ Understanding chemical bonding
 - ➡ Nanotechnology and biotechnology
 - ➡ ...
- Fundamental limitation: huge information space to explore.
- Fundamental advantage: huge data sets available.


Background: density functional theory (DFT)

- **Density functional theory** (DFT) is one of the most important computational method used to predict the electronic structure and properties of atoms, molecules, and solids. It is used to study the electronic structure of many-body quantum systems.
- Still too **slow**: $\sim 1\text{-}10\text{h}$ for a 20 atom molecule.
- This work approximates DFT results but its inference time is up to 300,000x faster!



Background: molecules as graphs


molecule



serotonin



graph



atoms \rightarrow nodes
bonds + distance \rightarrow edges

(graph is actually fully connected
because of spatial distance information)

Background: previous work on ML on graphs

- Two main branches of work on using machine learning on graphs:
 - ➔ **feature engineering**: Coulomb matrix, bag of bonds...
 - ➔ **neural networks**: relational networks, spectral networks, graph convolutional neural networks... These are invariant to graph isomorphism.
- This work aims to unify the different models of neural networks for graphs into one framework: **message passing neural networks** (MPNNs)
- This work **improves** on previous MPNNs algorithms by adding novel variations to this family of NNs.

Model: message passing neural network (MPNN)


- Consider an undirected graph G with node features x_v and edge features e_{vw} .
This graph has an associated adjacency matrix A .
- We aim to obtain some graph-level target.
- Computation is done in two phases.
 - (1) Message passing phase.
 - (2) Readout phase.
- The model requires learning three differentiable functions: message function M_t , vertex update functions U_t , and readout function R .

(1) Message passing phase

- Hidden states at each vertex are updated over T steps, indexed by t .
- This phase requires a:
 - ➔ Message passing function M_t .
 - ➔ Update function U_t .
- At each time-step $t \rightarrow t+1$ the hidden state on vector v is updated as follows:

$$(a) \quad m_v^{t+1} = \sum_{w \in N(v)} M_t(h_v^t, h_w^t, e_{vw})$$

$$(b) \quad h_v^{t+1} = U_t(h_v^t, m_v^{t+1})$$



(2) Readout phase

- The readout phase computes a feature vector y for the whole graph using some readout function R , such that:

$$\hat{y} = R(\{h_v^T \mid v \in G\})$$


- The function R must be invariant to permutation of the node states.

Model: previous MPNN work

- We can build different MPNNs by changing how we define the learnable functions M_t , U_t and R .
- Previous main work include:
 - ➡ Convolutional Networks for Learning Molecular Fingerprints, Duvenaud et al. (2015)
 - ➡ Gated Graph Neural Networks (GG-NN), Li et al. (2016)
 - ➡ Interaction Networks, Battaglia et al. (2016)
 - ➡ Molecular Graph Convolutions, Kearnes et al. (2016)
 - ➡ Deep Tensor Neural Networks, Schütt et al. (2017)

Model: novelties of this work

- Develop novelties around the GG-NN (Li et al.).

 : novel

- Functions:

(a) $M_t = A(e_{vw})h_w$

allows for vector valued edge features!

(b) $U_t = GRU(h_v^t, m_v^{t+1})$ same as in GG-NN

(c) $R = \text{set2set model}$

defined by Vinyals et al (2015)

- Adding **virtual edges** to link nodes that are not initially connected. Can even create a latent **master node** (connected to every input node).
- More **efficient propagation protocol**: break-up node embeddings.

Dataset and Task

- The dataset used in this problem is the **QM9** [1] which contains 130,000 organic molecules with at most 30 atoms of the elements H, C, O, N, F.
- There are 13 tasks that we will complete to compare performance with DFT.

Tasks include:

- ➡ Find the energy of the highest occupied molecular orbital.
- ➡ Heat capacity of the molecule
- ➡ ...

Code

Define message passing function:

```
def m_mpnn(self, h_v, h_w, e_vw, opt={}):
    # Matrices for each edge
    edge_output = self.learn_modules[0](e_vw)
    edge_output = edge_output.view(-1, self.args['out'], self.args['in'])

    h_w_rows = h_w[..., None].expand(h_w.size(0), h_v.size(1), h_w.size(1)).contiguous()

    h_w_rows = h_w_rows.view(-1, self.args['in'])

    h_multiply = torch.bmm(edge_output, torch.unsqueeze(h_w_rows, 2))

    m_new = torch.squeeze(h_multiply)

    return m_new

def out_mpnn(self, size_h, size_e, args):
    return self.args['out']

def init_mpnn(self, params):
    learn_args = []
    learn_modules = []
    args = {}

    args['in'] = params['in']
    args['out'] = params['out']

    # Define a parameter matrix A for each edge label.
    learn_modules.append(NNet(n_in=params['edge_feat'], n_out=(params['in']*params['out'])))

    return nn.ParameterList(learn_args), nn.ModuleList(learn_modules), args
```

Perform message passing:

```
m_t = {}
for v in g.nodes_iter():
    neigh = g.neighbors(v)
    m_neigh = type(h_t)
    for w in neigh:
        if (v,w) in e:
            e_vw = e[(v, w)]
        else:
            e_vw = e[(w, v)]
        m_v = m.forward(h_t[v], h_t[w], e_vw)
        if len(m_neigh):
            m_neigh += m_v
        else:
            m_neigh = m_v

    m_t[v] = m_neigh
```

Code

Define update function:

```
def u_ggnn(self, h_v, m_v, opt={}):
    h_v.contiguous()
    m_v.contiguous()
    h_new = self.learn_modules[0](torch.transpose(m_v, 0, 1), torch.unsqueeze(h_v, 0))[0]
    return torch.transpose(h_new, 0, 1)

def init_ggnn(self, params):
    learn_args = []
    learn_modules = []
    args = {}

    args['in_m'] = params['in_m']
    args['out'] = params['out']

    # GRU
    learn_modules.append(nn.GRU(params['in_m'], params['out']))

    return nn.ParameterList(learn_args), nn.ModuleList(learn_modules), args
```

Code

Take readout function from set2set (Vinyals et al)

Define final MPNN:

```
class MPNN(nn.Module):
    def __init__(self, in_n, hidden_state_size, message_size, n_layers, l_target, type='regression'):
        super(MPNN, self).__init__()

        # Define message
        self.m = nn.ModuleList(
            [MessageFunction('mpnn', args={'edge_feat': in_n[1], 'in': hidden_state_size, 'out': message_size})])

        # Define Update
        self.u = nn.ModuleList([UpdateFunction('mpnn',
                                                args={'in_m': message_size,
                                                      'out': hidden_state_size})])

        # Define Readout
        self.r = ReadoutFunction('mpnn',
                                  args={'in': hidden_state_size,
                                        'target': l_target})

        self.type = type

        self.args = {}
        self.args['out'] = hidden_state_size

        self.n_layers = n_layers
```

```
def forward(self, g, h_in, e):
    h = []

    # Padding to some larger dimension d
    h_t = torch.cat([h_in, Variable(
        torch.zeros(h_in.size(0), h_in.size(1), self.args['out'] - h_in.size(2)).type_as(h_in.data))], 2)

    h.append(h_t.clone())

    # Layer
    for t in range(0, self.n_layers):
        e_aux = e.view(-1, e.size(3))

        h_aux = h[t].view(-1, h[t].size(2))

        m = self.m[0].forward(h[t], h_aux, e_aux)
        m = m.view(h[0].size(0), h[0].size(1), -1, m.size(1))

        # Nodes without edge set message to 0
        m = torch.unsqueeze(g, 3).expand_as(m) * m

        m = torch.squeeze(torch.sum(m, 1))

        h_t = self.u[0].forward(h[t], m)

        # Delete virtual nodes
        h_t = (torch.sum(h_in, 2).expand_as(h_t) > 0).type_as(h_t) * h_t
        h.append(h_t)

    # Readout
    res = self.r.forward(h)

    if self.type == 'classification':
        res = nn.LogSoftmax()(res)

    return res
```

Results

- This work obtains **state-of-the-art** predictions of DFT on all 13 tasks.
- Metric is proportional to mean-squared error (the smaller the better)

	feature engineering + off-the-shelf classifiers					previous MPNN models			this work	
Target	BAML	BOB	CM	ECFP4	HDAD	GC[2]	GG-NN[3]	DTNN	enn-s2s	enn-s2s-ens5
mu	4.34	4.23	4.49	4.82	3.34	0.70	1.22	-	0.30	0.20
alpha	3.01	2.98	4.33	34.54	1.75	2.27	1.55	-	0.92	0.68
HOMO	2.20	2.20	3.09	2.89	1.54	1.18	1.17	-	0.99	0.74
LUMO	2.76	2.74	4.26	3.10	1.96	1.10	1.08	-	0.87	0.65
gap	3.28	3.41	5.32	3.86	2.49	1.78	1.70	-	1.60	1.23
R2	3.25	0.80	2.83	90.68	1.35	4.73	3.99	-	0.15	0.14
ZPVE	3.31	3.40	4.80	241.58	1.91	9.75	2.52	-	1.27	1.10
U0	1.21	1.43	2.98	85.01	0.58	3.02	0.83	-	0.45	0.33
U	1.22	1.44	2.99	85.59	0.59	3.16	0.86	-	0.45	0.34
H	1.22	1.44	2.99	86.21	0.59	3.19	0.81	-	0.39	0.30
G	1.20	1.42	2.97	78.36	0.59	2.95	0.78	.84 ²	0.44	0.34
Cv	1.64	1.83	2.36	30.29	0.88	1.45	1.19	-	0.80	0.62
Omega	0.27	0.35	1.32	1.47	0.34	0.32	0.53	-	0.19	0.15
Average	2.17	2.08	3.37	53.97	1.35	2.59	1.36	-	0.68	0.52


[2] Kearnes et al. *Molecular Graph Convolutions: Moving Beyond Fingerprints*. Journal of Computer-Aided Molecular Design, 2016.

[3] Li et al. *Gated Graph Sequence Neural Networks*. ICLR, 2016.

Future work

- Train on larger molecular spaces.
- Add attention mechanism to the incoming message vectors.
- Material science uses infinite periodic graphs.

molecule



molecool

